


Logistic regression



Michael Friendly
Psych 6136

<http://friendly.github.io/psy6136>



Model-based methods: Overview

Structure

- Explicitly assume some probability distribution for the data, e.g., binomial, Poisson, ...
- Distinguish between the **systematic** component— explained by the model— and a **random** component, which is not
- Allow a compact summary of the data in terms of a (hopefully) small number of parameters

Advantages

- Inferences: hypothesis tests *and* confidence intervals
- Can test **individual** model terms (**anova()**)
- Methods for model selection: adjust balance between goodness-of-fit and parsimony
- Predicted values give **model-smoothed** summaries for plotting
- \implies Interpret the fitted model graphically

loglm() vs. glm()

With **loglm()** you can only test overall fit (**anova()**) or difference between models (**Lrstats()**)

```
> berk.mod1 <- loglm(~ Dept * (Gender + Admit),  
data=UCBAdmissions)  
> berk.mod2 <- loglm(~(Admit + Dept + Gender)^2,  
data=UCBAdmissions)  
> anova(berk.mod2)  
Call:  
loglm(formula = ~(Admit + Dept + Gender)^2, data =  
UCBAdmissions)
```

Statistics:

	X^2	df	P(> X^2)
Likelihood Ratio	20.20	5	0.001144
Pearson	18.82	5	0.00207

What we can say:

Even the model with all pairwise associations fits poorly

Comparing models with `anova()` and `LRstats()`

```
> anova(berk.mod1, berk.mod2, test="Chisq")
LR tests for hierarchical log-linear models
```

Model 1:

```
~Dept * (Gender + Admit)
```

Model 2:

```
~(Admit + Dept + Gender)^2
```

	Deviance	df	Delta(Dev)	Delta(df)	P(> Delta(Dev))
Model 1	21.74	6			
Model 2	20.20	5	1.531	1	0.21593
Saturated	0.00	0	20.204	5	0.00114

```
> LRstats(berk.mod1, berk.mod2)
```

Likelihood summary table:

	AIC	BIC	LR	Chisq	Df	Pr(>Chisq)
berk.mod1	217	238		21.7	6	0.0014 **
berk.mod2	217	240		20.2	5	0.0011 **

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

loglm() vs. glm()

With **glm()** you can test individual terms using **anova()** or **car::Anova()**

```
> berkeley <- as.data.frame(UCBAdmissions)
> berk.glm2 <- glm(Freq ~ (Dept+Gender+Admit)^2, data=berkeley,
+                  family="poisson")
> anova(berk.glm2, test="Chisq")
Analysis of Deviance Table
```

```
Model: poisson, link: log
Response: Freq
```


Terms added sequentially (first to last)

	Df	Deviance	Resid.	Df	Resid.	Dev	Pr(>Chi)
NULL				23		2650	
Dept	5	160		18		2491	<2e-16 ***
Gender	1	163		17		2328	<2e-16 ***
Admit	1	230		16		2098	<2e-16 ***
Dept:Gender	5	1221		11		877	<2e-16 ***
Dept:Admit	5	855		6		22	<2e-16 ***
Gender:Admit	1	2		5		20	0.22

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Fitting & graphing models: Overview

Object-oriented approach in R:



- Fit model (`obj <- glm(...)`) → a model object
- `print(obj)` and `summary(obj)` → numerical results
- `anova(obj)` and `Anova(obj)` → tests for model terms
- `update(obj)`, `add1(obj)`, `drop1(obj)` for model selection

Plot methods:

- `plot(obj)` often gives diagnostic plots
- Other plot methods:
 - Mosaic plots: `mosaic(obj)` for "loglm" and "glm" objects
 - Effect plots: `plot(Effect(obj))` for nearly all linear models
 - Influence plots (car): `influencePlot(obj)` for "glm" objects

Objects & methods

How this works:

- Model objects have a "class" attribute:
 - `loglm()`: "loglm"
 - `glm()`: `c("glm", "lm")` — inherits also from `lm()`
- Class-specific methods have names like `method.class`, e.g., `plot.glm()`, `mosaic.loglm()`
- Generic functions (`print()`, `summary()`, `plot()` ...) call the appropriate method for the class

```
arth.mod <- glm(Better ~ Age + Sex + Treatment, data=Arthritis)
class(arth.mod)
```

```
## [1] "glm" "lm"
```

Objects & methods

Methods for “glm” objects

```
> library(MASS); library(vcdExtra)
> methods(class="glm")
[1] add1               addterm             anova               Anova
[5] asGnm              assoc               avPlot             avPlot3d
[9] Boot               bootCase            brief               ceresPlot
[13] coerce             confidenceEllipse   confint             Confint
[17] cooks.distance     deviance            drop1               dropterm
[21] effects            extractAIC           family              formula
[25] gamma.shape        influence            initialize           leveragePlot
[29] linearHypothesis   logLik              mcPlot              mmp
[33] model.frame        modFit              mosaic              ncvTest
[37] nobs               predict              print               profile
[41] qqPlot             residualPlot         residualPlots       residuals
[45] rootogram          rstandard           rstudent            S
[49] show               sieve               sigmaHat            slotsFromS3
[53] summary            vcov                weights
see '?methods' for accessing help and source code
```


There are many, many **plot()** methods for different types of objects
e.g., **plot()** for a “glm” object → **plot.glm()**

```
> methods("plot")
 [1] plot,ANY-method      plot,color-method    plot.acf*
 [4] plot.ca*             plot.correspondence* plot.data.frame*
 [7] plot.decomposed.ts*  plot.default         plot.dendrogram*
[10] plot.density*        plot.ecdf             plot.factor*
[13] plot.formula*        plot.function        plot.gnm*
[16] plot.goodfit*        plot.hcl_palettes*   plot.hclust*
[19] plot.histogram*     plot.HLtest*         plot.HoltWinters*
[22] plot.isoreg*         plot.lda*            plot.lm*
[25] plot.loddsratio*     plot.loglm*          plot.mca*
[28] plot.medpolish*      plot.mjca*           plot.mlm*
[31] plot.ppr*            plot.prcomp*         plot.princomp*
[34] plot.profile*        plot.profile.gnm*    plot.profile.nls*
[37] plot.qv*             plot.raster*         plot.ridgelm*
[40] plot.rootogram*     plot.shingle*        plot.spec*
[43] plot.stepfun         plot.stl*            plot.structable*
[46] plot.table*          plot.trellis*        plot.ts
[49] plot.tskernel*       plot.TukeyHSD*       plot.zoo*
see '?methods' for accessing help and source code
```

Modeling approaches: Overview

Association models

- Loglinear models
(contingency table form)
[Admit][Gender Dept]
[Admit Dept][Gender Dept]
[AdmitDept][AdmitGender][GenderDept]
- Poisson GLMs
(Frequency data frame)
 $\text{Freq} \sim \text{Admit} + \text{Gender} * \text{Dept}$
 $\text{Freq} \sim \text{Admit} * \text{Dept} + \text{Gender} * \text{Dept}$
 $\text{Freq} \sim \text{Admit} * (\text{Dept} + \text{Gender}) + \text{Gender} * \text{Dept}$
- Ordinal variables
 $\text{Freq} \sim \text{right} + \text{left} + \text{Diag}(\text{right}, \text{left})$
 $\text{Freq} \sim \text{right} + \text{left} + \text{Symm}(\text{right}, \text{left})$

Response models

- Binary response
- Categorical predictors: logit models
 $\text{logit}(\text{Admit}) \sim 1$
 $\text{logit}(\text{Admit}) \sim \text{Dept}$
 $\text{logit}(\text{Admit}) \sim \text{Dept} + \text{Gender}$
- Continuous/mixed predictors
- Logistic regression models
 $\text{Pr}(\text{Admit}) \sim \text{Dept} + \text{Gender} + \text{Age} + \text{GRE}$
- Polytomous response
- Ordinal: proportional odds model
 $\text{Improve} \sim \text{Age} + \text{Sex} + \text{Treatment}$
- General multinomial model
 $\text{WomenWork} \sim \text{Kids} + \text{HusbandIncome}$

Logistic regression

Response variable


- Binary response: success/failure, vote: yes/no
- Binomial data: x successes in n trials (grouped data)
- Ordinal response: none < some < severe depression
- Polytomous response: vote Liberal, Tory, NDP, Green

Explanatory variables

- Quantitative regressors: age, dose
- Transformed regressors: $\sqrt{\text{age}}$, $\log(\text{dose})$
- Polynomial regressors: age^2 , age^3 , \dots (or better: splines)
- Categorical predictors: treatment, sex (dummy variables, contrasts)
- Interaction regressors: treatment \times age, sex \times age


This is exactly the same as in classical ANOVA, regression models

Example: Arthritis treatment



- The response variable, `Improved` is ordinal: "None" < "Some" < "Marked"
- A binary logistic model can consider just `Better = (Improved > "None")`
- Other important predictors: Sex, Treatment
- Main Q: how does treatment affect outcome?
- How does this vary with Age and Sex?
- This plot shows the binary observations, with several model-based smoothings


Example: Berkeley admissions



- Admit/Reject can be considered a **binomial response** for each Dept and Gender
- Logistic regression here is analogous to an ANOVA model, but for log odds(Admit)
- (With categorical predictors, these are often called **logit** models)
- Every such model has an equivalent **loglinear** model form.
- This plot shows fitted logits for the main effects model, $\text{Dept} + \text{Gender}$


Example: Survival in the Donner party

- Binary response: **survived**
- Categorical predictors: **sex**, **family**
- Quantitative predictor: **age**
- Q: Is the effect of age linear?
- Q: Are there interactions among predictors?
- This is a **generalized pairs plot**, with different plots for each pair



Binary response: What's wrong with OLS?

- For a binary response, $Y \in (0, 1)$, want to predict $\pi = \Pr(Y = 1 | x)$
- A **linear probability model** uses classical linear regression (OLS)
- Problems:
 - Gives predicted values and CIs outside $0 \leq \pi \leq 1$
 - Homogeneity of variance is violated: $\mathcal{V}(\hat{\pi}) = \hat{\pi}(1 - \hat{\pi}) \neq \text{constant}$
 - Inferences, hypothesis tests are wrong!



Linear regression vs Logistic regression

OLS regression:

- Assume $y|x \sim N(0, \sigma^2)$




Fig. 2.1. Graphical representation of a simple linear normal regression.

Logistic regression:

- Assume $\Pr(y=1|x) \sim \text{binomial}(p)$





Fig. 2.2. Graphical representation of a simple linear logistic regression.

y linear with x
constant residual variance

$y \sim \text{logit}(x)$
non-constant residual variance $\sim p(1-p)$

Logistic regression

- Logistic regression avoids these problems
- Models $\text{logit}(\pi_i) \equiv \log[\pi/(1 - \pi)]$
- logit is interpretable as “log odds” that $Y = 1$
- A related **probit** model gives very similar results, but is less interpretable
- For $0.2 \leq \pi \leq 0.8$ fitted values are close to those from linear regression.



Logistic regression: One predictor

For a single quantitative predictor, x , the simple **linear logistic regression model** posits a linear relation between the **log odds** (or **logit**) of $\Pr(Y = 1)$ and x ,

$$\text{logit}[\pi(x)] \equiv \log \left(\frac{\pi(x)}{1 - \pi(x)} \right) = \alpha + \beta x .$$

- When $\beta > 0$, $\pi(x)$ and the log odds increase as x increases; when $\beta < 0$ they decrease with x .
- This model can also be expressed as a model for the probabilities $\pi(x)$

$$\pi(x) = \text{logit}^{-1}[\pi(x)] = \frac{1}{1 + \exp[-(\alpha + \beta x)]}$$

Thinking logistically:

- Model is for the **log odds** of the marked response, $Y = 1$
- Can always back transform with logit^{-1} to get **probability** of $Y = 1$

Logistic regression: One predictor

The coefficients, α , β of this model have simple interpretations in terms of odds & log odds

$$\text{logit}[\pi(x)] \equiv \log\left(\frac{\pi(x)}{1 - \pi(x)}\right) = \alpha + \beta x \quad \text{odds}(Y = 1) \equiv \frac{\pi(x)}{1 - \pi(x)} = \exp(\alpha + \beta x) = e^\alpha (e^\beta)^x$$

β is the change in log odds for a unit increase in x

→ The odds of $Y=1$ are multiplied by e^β for each unit increase in x

α is the log odds when $x=0$

→ The odds of $Y=1$ when $x=0$ is e^α

In R, use `exp(coef(model))` to get these values

Another interpretation: In terms of probability, the slope of the logistic regression curve is $\beta\pi(1-\pi)$

This has the maximum value $\beta/4$ when $\pi = 1/2$

Logistic regression: Multiple predictors

- For a binary response, $Y \in (0, 1)$, let \mathbf{x} be a vector of p regressors, and π_i be the probability, $\Pr(Y = 1 \mid \mathbf{x})$.
- The logistic regression model is a linear model for the *log odds*, or *logit* that $Y = 1$, given the values in \mathbf{x} ,

$$\begin{aligned}\text{logit}(\pi_i) \equiv \log \left(\frac{\pi_i}{1 - \pi_i} \right) &= \alpha + \mathbf{x}_i^T \boldsymbol{\beta} \\ &= \alpha + \beta_1 x_{i1} + \beta_2 x_{i2} + \cdots + \beta_p x_{ip}\end{aligned}$$

- An equivalent (non-linear) form of the model may be specified for the probability, π_i , itself,

$$\pi_i = \{1 + \exp(-[\alpha + \mathbf{x}_i^T \boldsymbol{\beta}])\}^{-1}$$

- The logistic model is also a *multiplicative* model for the odds of “success,”

$$\frac{\pi_i}{1 - \pi_i} = \exp(\alpha + \mathbf{x}_i^T \boldsymbol{\beta}) = \exp(\alpha) \exp(\mathbf{x}_i^T \boldsymbol{\beta})$$

Increasing x_{ij} by 1 increases $\text{logit}(\pi_i)$ by β_j , and multiplies the odds by e^{β_j} .

Fitting the logistic regression model

Logistic regression models are the special case of generalized linear models, fit in R using `glm(..., family=binomial)`

For this example, we define **Better** as any improvement at all

```
> data(Arthritis, package="vcd")
> Arthritis$Better <- as.numeric(Arthritis$Improved > "None")
```

Fit and print:

```
> (arth.logistic <- glm(Better ~ Age, data=Arthritis, family=binomial))

Call:  glm(formula = Better ~ Age, family = binomial, data = Arthritis)

Coefficients:
(Intercept)          Age
    -2.6421         0.0492

Degrees of Freedom: 83 Total (i.e. Null);  82 Residual
Null Deviance:      116
Residual Deviance: 109    AIC: 113
```

The summary() method gives details and tests of coefficients

```
> summary(arth.logistic)
```

Call:

```
glm(formula = Better ~ Age, family = binomial, data = Arthritis)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-1.5106	-1.1277	0.0794	1.0677	1.7611

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-2.6421	1.0732	-2.46	0.014 *
Age	0.0492	0.0194	2.54	0.011 *

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance:	116.45	on 83	degrees of freedom
Residual deviance:	109.16	on 82	degrees of freedom

Interpreting coefficients

```
> coef(arth.logistic)
(Intercept)    Age
-2.64207    0.04925
```

```
> exp(coef(arth.logistic))
(Intercept)    Age
0.07121    1.05048
> exp(10*coef(arth.logistic)[2])
Age
1.636
```

Interpretations:

- log odds(Better) increase by $\beta = 0.0492$ for each year of age
- odds(Better) multiplied by $e^{\beta} = 1.05$ for each year of age— a 5% increase
- over 10 years, odds(Better) are multiplied by $\exp(10 \times 0.0492) = 1.64$, a 64% increase.
- $\Pr(\text{Better})$ increases by $\beta/4 = 0.0123$ for each year (near $\pi = \frac{1}{2}$)

Multiple predictors

The main interest here is the effect of Treatment. Sex and Age are **control variables**. Fit the **main effects** model (no interactions):

$$\text{logit}(\pi_i) = \alpha + \beta_1 x_{i1} + \beta_2 x_{i2} + \beta_3 x_{i3}$$

where x_1 is *Age* and x_2 and x_3 are the factors representing *Sex* and *Treatment*, respectively. R uses dummy (0/1) variables for factors.

$$x_2 = \begin{cases} 0 & \text{if Female} \\ 1 & \text{if Male} \end{cases} \quad x_3 = \begin{cases} 0 & \text{if Placebo} \\ 1 & \text{if Treatment} \end{cases}$$

- α doesn't have a sensible interpretation here. Why?
- β_1 : increment in log odds(Better) for each year of age.
- β_2 : difference in log odds for male as compared to female.
- β_3 : difference in log odds for treated vs. the placebo group

Multiple predictors: Fitting

Fit the main effects model. Use $I(\text{Age} - 50)$ to center Age, making α interpretable

```
arth.logistic2 <- glm(Better ~ I(Age - 50) + Sex + Treatment,  
                      data=Arthritis, family=binomial)
```

lmtest::coeftest() gives just the tests of coefficients provided by summary()

```
> lmtest::coeftest(arth.logistic2)
```

z test of coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-0.5781	0.3674	-1.57	0.116
I(Age - 50)	0.0487	0.0207	2.36	0.018 *
SexMale	-1.4878	0.5948	-2.50	0.012 *
TreatmentTreated	1.7598	0.5365	3.28	0.001 **

broom::glance() gives model fit statistics

```
> broom::glance(arth.logistic2)
```

```
# A tibble: 1 x 8
```

	null.deviance	df.null	logLik	AIC	BIC	deviance	df.residual	nobs
	<dbl>	<int>	<dbl>	<dbl>	<dbl>	<dbl>	<int>	<int>
1	116.	83	-46.0	100.	110.	92.1	80	84

Interpreting coefficients

```
> cbind(coef=coef(arth.logistic2),  
+       OddsRatio=exp(coef(arth.logistic2)),  
exp(confint(arth.logistic2)))  
Waiting for profiling to be done...
```

	coef	OddsRatio	2.5 %	97.5 %
(Intercept)	-0.5781	0.561	0.2647	1.132
I(Age - 50)	0.0487	1.050	1.0100	1.096
SexMale	-1.4878	0.226	0.0652	0.689
TreatmentTreated	1.7598	5.811	2.1187	17.727

- $\alpha = -0.578$: At age 50, females given placebo have odds(Better) of $e^{-0.578} = 0.56$.
- $\beta_1 = 0.0487$: Each year of age multiplies odds(Better) by $e^{0.0487} = 1.05$, a 5% increase.
- $\beta_2 = -1.49$: Males $e^{-1.49} = 0.26 \times$ less likely to show improvement as females. (Or, females $e^{1.49} = 4.437 \times$ more likely than males.)
- $\beta_3 = 1.76$: Treated $e^{1.76} = 5.81 \times$ more likely Better than Placebo

Hypothesis testing: Questions

- **Overall test:** How does my model, $\text{logit}(\pi) = \alpha + \mathbf{x}^T \beta$ compare with the null model, $\text{logit}(\pi) = \alpha$?

$$H_0 : \beta_1 = \beta_2 = \cdots = \beta_p = 0$$

- **One predictor:** Does x_k significantly improve my model? Can it be dropped?

$$H_0 : \beta_k = 0 \quad \text{given other predictors retained}$$

- **Lack of fit:** How does my model compare with a perfect model (saturated model)?

For ANOVA, regression, these tests are carried out using F -tests and t -tests. In logistic regression (fit by maximum likelihood) we use

- F -tests \rightarrow likelihood ratio G^2 tests
- t -tests \rightarrow Wald z or χ^2 tests

Maximum likelihood estimation

In classical linear models using `lm()`, we fit using ordinary least squares.
All `glm()` models use maximum likelihood estimation— better properties

- Likelihood, $\mathcal{L} = \Pr(\text{data} \mid \text{model})$, as function of model parameters
- For case i ,

$$\mathcal{L}_i = \begin{cases} p_i & \text{if } Y = 1 \\ 1 - p_i & \text{if } Y = 0 \end{cases} = p_i^{Y_i} (1 - p_i^{Y_i}) \quad \text{where} \quad p_i = 1 / (1 + \exp(\mathbf{x}_i \boldsymbol{\beta}))$$

- Under independence, joint likelihood is the product over all cases

$$\mathcal{L} = \prod_i^n p_i^{Y_i} (1 - p_i^{Y_i})$$

- \implies Find estimates $\hat{\boldsymbol{\beta}}$ that maximize $\log \mathcal{L}$. Iterative, but this solves the “estimating equations”

$$\mathbf{X}^T \mathbf{y} = \mathbf{X}^T \hat{\mathbf{p}}$$

Overall model tests

Likelihood ratio test (G^2)

- Compare **nested** models, similar to F tests in OLS
- Let L_1 = maximized value for **our model**

$$\text{logit}(\pi_i) = \beta_0 + \mathbf{x}_i^T \boldsymbol{\beta} \quad \text{w/ } k \text{ predictors}$$

- Let L_0 = maximized likelihood for the **null model**

$$\text{logit}(\pi_i) = \beta_0 \quad \text{under } H_0: \beta_1 = \beta_2 = \dots = \beta_k$$

- Likelihood ratio test statistic:

$$G^2 = -2 \log \left(\frac{L_0}{L_1} \right) = 2(\log L_1 - \log L_0) \sim \chi_k^2$$

Wald tests & confidence intervals

- Analogous to t -tests in OLS

- Test $H_0: \beta_i = 0$

$$z = \frac{b_i}{s(b_i)} \sim \mathcal{N}(0,1) \quad \text{or} \quad z^2 \sim \chi_1^2$$

- Confidence interval

$$b_i \pm z_{1-\alpha/2} s(b_i)$$

```
> r1 <- lmtest::coeftest(arth.logistic2)
```

```
> r2 <- confint(arth.logistic2)
```

```
Waiting for profiling to be done...
```

```
> cbind(r1, r2)
```


	Estimate	Std. Error	z value	Pr(> z)	2.5 %	97.5 %
(Intercept)	-0.578	0.367	-1.6	0.116	-1.33	0.124
I(Age - 50)	0.049	0.021	2.4	0.018	0.01	0.092
SexMale	-1.488	0.595	-2.5	0.012	-2.73	-0.372
TreatmentTreated	1.760	0.536	3.3	0.001	0.75	2.875

LR, Wald & Score tests

Testing Global Null Hypothesis: BETA=0

$H_0: \beta_1 = \beta_2 = \beta_3 = 0$

Test	Chi-Square	DF	Pr > ChiSq
Likelihood Ratio	24.3859	3	<.0001
Score	22.0051	3	<.0001
Wald	17.5147	3	0.0006




Different ways to measure departure from $H_0: \beta = 0$

- LR test: diff^{ce} in $\log L$
- Wald test: $(\hat{\beta} - \beta_0)^2$
- Score test: slope at $\beta = 0$

Plotting logistic regression data


Plotting a binary response together with a fitted logistic model can be difficult because the 0/1 response leads to much overplotting.

- Need to jitter the points
- Useful to show the fitted logistic curve
- Confidence band gives a sense of uncertainty
- Adding a non-parametric (loess) smooth shows possible nonlinearity
- NB: Can plot either on the **response** scale (probability) or the **link** scale (logit) where effects are linear




Types of plots

- **Conditional plots:** Stratified plot of Y or $\text{logit}(Y)$ vs. one X , conditioned by other predictors--- only that subset is plotted for each






Types of plots

- **Full-model plots:** Plot of fitted response surface, showing all effects; usually shown in several panels



Types of plots


- **Effect plots:** plots of predicted effects for terms in the model, averaged over predictors not shown in a given plot



Conditional plots with ggplot2

Plot Arthritis data by Treatment, ignoring Sex; overlay fitted logistic reg. lines

```
gg <- ggplot(Arthritis, aes(Age, Better, color=Treatment)) +  
  xlim(15, 85) +  
  geom_jitter(height = 0.02, width = 0, size=2) +  
  stat_smooth(method = "glm", family = binomial, alpha = 0.2,  
             aes(fill=Treatment), size=2.5, fullrange=TRUE) +  
  theme_bw(base_size = 16) + theme(legend.position = c(.85, .2))  
gg # show the plot
```




`geom_jitter()` shows the observations more distinctly

Fitted lines use `method="glm"`, `family=binomial`

Conditional plots with ggplot2

Can show the conditional plots for M & F, simply by faceting by Sex

```
gg + facet_wrap(~ Sex)
```



Only the data for each Sex is used in each plot

Plotting the data points shows that the data for males is too thin to give good estimates of separate regression

Full-model plots

Full-model plots show the fitted values on the logit scale or on the response scale (probability), usually with confidence bands. This often requires a bit of custom programming.

Steps:

- Obtain fitted values with `predict(model, se.fit=TRUE)` — `type="link"` (logit) is the default
- Can use `type="response"` for probability scale
- Join this to your data (`cbind()`)
- Plot as you like: `plot()`, `ggplot()`, ...

```
> arth.fit2 <- cbind(Arthritis,  
+                   predict(arth.logistic2, se.fit = TRUE))  
> head(arth.fit2[, -9], 4)
```

	ID	Treatment	Sex	Age	Improved	Better	fit	se.fit
1	57	Treated	Male	27	Some	1	-1.43	0.758
2	46	Treated	Male	29	None	0	-1.33	0.728
3	77	Treated	Male	30	None	0	-1.28	0.713
4	17	Treated	Male	32	Marked	1	-1.18	0.684

Plotting with ggplot2

Plot the fitted log odds, confidence band and observations


```
arth.fit2 <- arth.fit2 |>
  mutate(obs = ifelse(Better==0, -4, 4))    # show obs at -4, 4

ggplot( arth.fit2, aes(x=Age, y=fit, color=Treatment)) +
  geom_line(size = 2) +
  geom_ribbon(aes(ymin = fit - 1.96 * se.fit,
                 ymax = fit + 1.96 * se.fit,
                 fill = Treatment), alpha = 0.2,
            color = "transparent") +
  labs(x = "Age", y = "Log odds (Better)") +
  geom_jitter(aes(y=obs), height=0.25, width=0) +
  facet_wrap(~ Sex) +
  theme_bw(base_size = 16)
```

Using color=Treatment gives separate points and lines for the two groups

Full-model plot

Plotting on the logit scale shows the **additive** effects of age, treatment and sex
NB: easier to compare the treatment groups within the **same** panel




These plots show model uncertainty (confidence bands)
Jittered points show the data

Full-model plot

Plotting on the probability scale may be simpler to interpret

Use `predict(... type = "response")` to get fitted probabilities

```
arth.fit2r <- cbind(Arthritis,  
                    predict(arth.logistic2, se.fit = TRUE, type="response"))
```



Models with interactions

Is the linear effect of age the same for females, males?

- We can test this by adding an **interaction** of Sex x Age
- **update()** makes it easy to add/subtract terms from a model
- **car::Anova()** gives partial tests of each term after all others

```
> arth.logistic4 <- update(arth.logistic2, . ~ . + I(Age-50):Sex)
> car::Anova(arth.logistic4)
Analysis of Deviance Table (Type II tests)
```


Response: Better

	LR	Chisq	Df	Pr(>Chisq)
I(Age - 50)	6.16	1		0.01308 *
Sex	6.98	1		0.00823 **
Treatment	11.90	1		0.00056 ***
I(Age - 50):Sex	3.42	1		0.06430 .

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

The interaction term Age:Sex is not quite significant, but plot the fitted model anyway

Models with interactions



- Only the model changes
- `predict()` automatically incorporates the revised model terms
- Plotting steps remain the same
- This interpretation is quite different!

Effect plots: Basic ideas

Show a given **marginal** effect, **controlling** / adjusting for other model effects

Data

	x1	x2	sex	x1x2	y	yhat
1	1	1	F	1	4.73	4.46
2	2	1	M	0	6.10	5.55
3	3	1	F	-1	4.32	4.34
4	1	1	F	1	4.84	4.46
5	2	1	F	0	4.73	4.40
...
29	2	2	M	0	6.10	6.15
30	3	2	F	1	6.71	7.14


• Fit data: $\mathbf{X}\hat{\boldsymbol{\beta}} \Rightarrow \hat{\mathbf{y}}$


• Score data $\mathbf{X}^* \hat{\boldsymbol{\beta}} \Rightarrow \hat{\mathbf{y}}^*$

- plot vars: vary over range
- control vars: fix at means


Score data

	x1	x2	sex	x1:x2	y	yhat *
31	1	1	0.5	1	NA	5.030
32	2	1	0.5	2	NA	4.971
33	3	1	0.5	3	NA	4.912
34	1	2	0.5	2	NA	3.437
35	2	2	0.5	4	NA	5.574
36	3	2	0.5	6	NA	7.710


 plot vars


 control vars

plot



Effect plots: Details


- For simple models, full model plots show the complete relation between response and *all predictors*.
- Fox(1987)— For complex models, often wish to plot a specific main effect or interaction (including lower-order relatives)— *controlling for other effects*
 - Fit full model to data with linear predictor (e.g., logit) $\eta = \mathbf{X}\beta$ and link function $g(\mu) = \eta \rightarrow$ estimate \mathbf{b} of β and covariance matrix $\widehat{V}(\mathbf{b})$ of \mathbf{b} .
 - Construct “score data”
 - Vary each predictor in the term over its’ range
 - Fix other predictors at “typical” values (mean, median, proportion in the data)
 - \rightarrow “effect model matrix,” \mathbf{X}^*
 - Use `predict()` on \mathbf{X}^*
 - Calculate fitted effect values, $\hat{\eta}^* = \mathbf{X}^* \mathbf{b}$.
 - Standard errors are square roots of $\text{diag } \mathbf{X}^* \widehat{V}(\mathbf{b}) \mathbf{X}^{*\top}$
 - Plot $\hat{\eta}^*$, or values transformed back to scale of response, $g^{-1}(\hat{\eta}^*)$.
- *Note:* This provides a general means to visualize interactions in *all* linear and generalized linear models.

Plotting main effects


allEffects() calculates effects for all high-order terms in the model
The response is plotted on the logit scale, but labeled with probabilities

```
library(effects)  
arth.eff2 <- allEffects(arth.logistic2)  
plot(arth.eff2, rows=1, cols=3, lwd=2)
```


Age effect plot



Sex effect plot



Treatment effect plot



Averaged
over:

Sex
Treatment


Age
Treatment

Age
Sex

Full-model plot



The full-model plot is simply the **Effect()** of the highest-order interaction of factors

```
arth.full <- Effect(c("Age", "Treatment", "Sex"), arth.logistic2)
plot(arth.full, multiline=TRUE, ci.style="bands",
     colors = c("red", "blue"), lwd=3, . . .)
```



Model with interaction of Age × Sex

```
arth.eff4 <- allEffects(arth.logistic4)  
plot(arth.eff4, lwd=2)
```



Only the high-order terms: Treatment & Age * Sex are shown & need to be interpreted

Q: How would you describe this?

Race & Crime

Toronto Star investigation of racial disparities in treatment by Toronto Police Services

FOI request → > ½ M arrests, 1997—2002

Evidence for racial profiling?

Only look at discretionary charges:

Simple marijuana possession
Non-moving auto infractions

THE SATURDAY STAR

The photo that never was

GARTH WOOLSEY, C3



Periods of rain; windy. High 14 C

October 19, 2002

thestar.com ONTARIO EDITION

Also inside ...

- **Waterfront:** Dreams of what could be, B1, B4-5
- **Hydro woes:** Insulating against price spikes, E1
- **Wheels:** The Bug goes roofless, G1
- **Paul Martin:** The man who would be king, H1
- **Carol Shields:** Her last novel? Unless ... J1

AN INVESTIGATION INTO RACE AND CRIME




SUING POLICE: Jason Burke, falsely accused of dealing drugs during Caribana two years ago, says he was a victim of racial profiling.

Singled out

Star analysis of police crime data shows justice is different for blacks and whites

Telling numbers

Police records show that a black person in Toronto arrested on a single drug possession charge was less likely to be released at the scene...



Blacks arrested by Toronto police are treated more harshly than whites, a Toronto Star analysis of crime data shows.

Black people, charged with simple drug possession, are taken to police stations more often than whites facing the same charge.

Once at the station, accused blacks are held overnight, for a bail hearing, at twice the rate of whites.

The Toronto crime data also shows a disproportionate number of black motorists are ticketed for violations that only surface following a traffic stop. This difference, say civil libertarians, community

Managing Editor's notebook, A2

leaders and criminologists, suggests police use racial profiling in deciding whom to pull over.

The evidence is contained in a massive police database recording more than 480,000 incidents in which an individual was arrested, or ticketed, for an offence dating back to 1996. It included almost 800,000 criminal and other charges. The Star obtained that data through a freedom of information request, marking the first time access to these numbers was granted to anyone outside the police

community.

Police are forbidden, by their governing board, from analyzing this data in terms of race, but The Star has no such restriction. The findings provide hard evidence of what blacks have long suspected — race matters in Canadian society especially when dealing with police.

Chief Julian Fantino disputed the findings, saying the colour of a person's skin has nothing to do with how they're treated by his officers.

"We don't treat people different-

Chrétien expected to keep cabinet minister

Ethics report has 'wigggle room' to save MacAulay

BY TIM HARPER
AND LES WHITTINGTON
OTTAWA BUREAU

OTTAWA — Jean Chrétien receives a report from his ethics counsellor today that is expected to give him enough "wigggle room" to keep his solicitor-general, Lawrence MacAulay, in the federal cabinet.

Ethics counsellor Howard Wilson completed his report and delivered it to the Prime Minister's Office last night, where it was received by Chrétien's chief of staff, Percy Downie.

It was then to be relayed to Chrétien by secure fax to Beirut, where the Prime Minister is attending a summit of French-speaking nations. It was 1:30 a.m. in Beirut when the fax arrived so Chrétien would likely be reading it this morning.

Senior sources said last night that unless there is a surprise in Wilson's report, the Prince Edward Island minister will remain, Chrétien will return to Ottawa and weather the inevitable storm of opposition and media protest and forge ahead with an ethics package by mid-week.

Wilson has been investigating whether MacAulay broke ethics guidelines for cabinet ministers in the awarding of a contract and extension worth \$100,000 to Everett Roche, a Charlottetown political friend of the solicitor-general's.

Chrétien will not fire MacAulay unless he is given incontrovertible evidence of wrongdoing for two key reasons, source-

■ Please see MacAulay, A8

INSIDE

Barclay L2
Births B7

Ellie Teshler L2
James Travers H2

Case study: Arrests for marijuana

- In Dec. 2002, the *Toronto Star* examined the issue of **racial profiling**, by analyzing a data base of 600,000+ arrest records from 1997-2002.
- They focused on a subset of arrests for which police action was **discretionary**, e.g., simple possession of small quantities of marijuana, where the police could:
 - Release the arrestee with a summons – like a parking ticket
 - Bring to police station, hold for bail, ... -- harsher treatment
- Response variable: released: “Yes”, “No”
 - Main predictor of interest: skin-colour of arrestee (black, white)
 - Other predictors: year, age, sex, ...

Toronto Star meets mosaic displays


How to communicate these results most effectively?

- What is the message? What features are directly comprehensible to the audience?

B SECTION » TORONTO STAR « WEDNESDAY, DECEMBER 11, 2002 ★ thestar.com

Race and Crime

Graphic designer's
early attempts



My early
attempts

York University professor Michael Friendly's expert statistical analysis provided confirmation for the Toronto Star's series on racial profiling by city police.

Man behind the numbers

Racial profiling: Presentation graphic

Together, we created this (nearly) **self-explaining** infographic

Title gives the main conclusion

Text description gives details

Bar width ~ charges
Divided by % release

numbers shown in the cells

Legend gives a layman's description of shading levels

Same charge, different treatment

Statistical analysis of single drug possession charges shows that blacks are much less likely to be released at the scene and much more likely to be held in custody for a bail hearing. Darker colours represent a stronger statistical link between skin colour and police treatment.

Degree of likelihood

- Much less* likely to occur
- Much more* likely to occur
- More likely* to occur


Whites are more likely to be released at the scene

6,662 charges laid



Blacks are much more likely to be held for bail hearings

2,446 charges laid



SOURCE: Toronto police arrest records 1996-2002

Arrests for marijuana: Data

Response variable: released

Control variables:

- year, age, sex
- employed, citizen: Yes, No
- checks: # of police databases (previous arrests, convictions, parole status) where the arrestee's name was found

```
> library(car)          # for Anova()
> data(Arrests, package = "carData")
> some(Arrests)
```

	released	colour	year	age	sex	employed	citizen	checks
218	Yes	White	2000	24	Male	Yes	Yes	0
1301	No	Black	1999	17	Male	Yes	No	1
1495	Yes	White	1998	23	Male	Yes	Yes	0
1732	Yes	Black	2000	18	Male	Yes	Yes	2
1838	Yes	Black	1997	27	Male	No	Yes	5
2257	No	White	2001	19	Male	No	Yes	2
3100	No	Black	2000	19	Male	No	Yes	4
3843	Yes	White	1999	20	Male	Yes	Yes	0
4580	Yes	Black	1999	26	Male	Yes	Yes	1
4833	Yes	Black	1998	38	Male	Yes	Yes	0

Arrests for marijuana: Model

year is numerical. But may be non-linear. Convert to a **factor**

Fit model with all main effects, but allow **interactions** of colour:year and colour:age

```
> Arrests$year <- as.factor(Arrests$year)
> arrests.mod <- glm(released ~ employed + citizen + checks +
                     colour*year + colour*age,
                     family=binomial, data=Arrests)
> Anova(arrests.mod)
```

Analysis of Deviance Table (Type II tests)


Response: released

	LR	Chisq	Df	Pr(>Chisq)
employed	72.7	1	< 2e-16	***
citizen	25.8	1	3.8e-07	***
checks	205.2	1	< 2e-16	***
colour	19.6	1	9.7e-06	***
year	6.1	5	0.29785	
age	0.5	1	0.49827	
colour:year	21.7	5	0.00059	***
colour:age	13.9	1	0.00019	***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Effect plot: Skin colour

```
plot(Effect("colour", arrests.mod), lwd=3, ci.style="bands", ...)
```




- Effect plot for colour shows average effect **controlling** (adjusting) for **all** other factors simultaneously
- (The *Star* analysis controlled for these one at a time.)
- Evidence for different treatment of blacks & whites
- Even Francis Nunziata could understand this.
- However, effect smaller than reported by the *Star*

Effect plots: Interactions

The story turned out to be more nuanced than reported by the *Toronto Star*

```
plot(Effect(c("colour", "year"), arrests.mod), multiline=TRUE, ...)
```



Up to 2000, strong evidence for differential treatment of blacks & whites


Also: evidence to support Police claim of effect of training to reduce racial effects in treatment

Effect plots: Interactions

A more surprising finding ...

```
plot(Effect(c("colour", "year"), arrests.mod), multiline=TRUE, ...)
```

Effects of skin colour and age on release



Opposite age effects for blacks & whites:


- Young blacks treated **more** harshly than young whites
- Older blacks treated **less** harshly than older whites

Effect plots: allEffects


All high-order terms can be viewed together using `plot(allEffects(mod))`

```
arrests.effects <- allEffects(arrests.mod,  
xlevels=list(age=seq(15,45,5)))  
plot(arrests.effects, ylab="Probability(released)", ...)
```


employed effect plot




citizen effect plot




checks effect plot



colour*year effect plot



colour*age effect plot



Model diagnostics

As in regression and ANOVA, the validity of a logistic regression model is threatened when:

- Important predictors have been omitted from the model
- Predictors assumed to be **linear** have **non-linear** effects on $\Pr(Y = 1)$
- Important **interactions** have been omitted
- A few “wild” observations have a large impact on the fitted model or coefficients

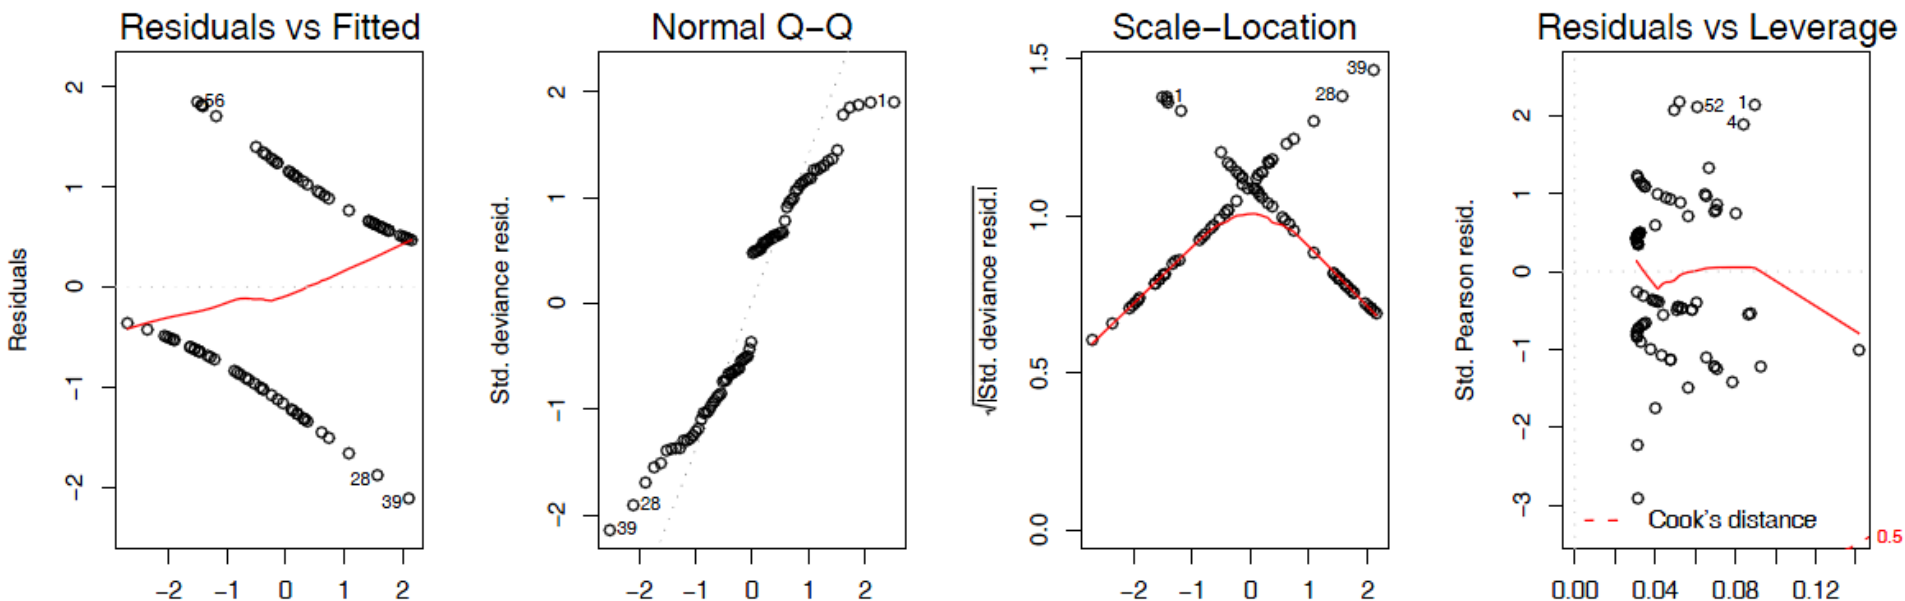
Model specification: Tools and techniques

- Use non-parametric smoothed curves to detect non-linearity
- Consider using polynomial terms (X^2, X^3, \dots) or **regression splines** (e.g., `ns(X, 3)`)
- Use `update(model, ...)` to test for interactions— formula: $. \sim .^2$

Diagnostic plots in R

In R, plotting a `glm` object gives the “regression quartet” – 4 basic diagnostic plots

```
arth.mod1 <- glm(Better ~ Age + Sex + Treatment, data=Arthritis,  
                 family='binomial')  
plot(arth.mod1)
```






These plots often look peculiar for logistic regression models
Better versions are available in the `car` package

Unusual data: Leverage & Influence

- “Unusual” observations can have dramatic effects on least-squares estimates in linear models
- Three archetypal cases:
 - Typical X (low leverage), bad fit -- Not much harm
 - Unusual X (high leverage), good fit -- Not much harm
 - Unusual X (high leverage), bad fit -- **BAD, BAD, BAD**
- Influential observations: unusual in *both* X & Y
- Heuristic formula:

$$\text{Influence} = X \text{ leverage} \times Y \text{ residual}$$



Influence plots

Influence (Cook's D) measures impact of individual obs. on coefficients, fitted values


$$\text{Influence} \sim \text{Residual } (y - \hat{y}) \times \text{Hat-value } (\mathbf{X} - \bar{\mathbf{X}})^2$$

Bubble size \sim influence

`influencePlot(mod1)`


Bad fit

Bad fit



Influence plots in R

```
library(car)  
influencePlot(arth.logistic2, ...)
```




X axis: Leverage ("hat values")
notable values: $> 2k/n$, $3k/n$

Y axis: Studentized residuals

Bubble size \sim Cook's D
(influence on coefficients)

Which cases are influential?

	Treatment	Sex	Age	Better	StudRes	Hat	CookD
1	Treated	Male	27	1	1.92	0.0897	0.1128
4	Treated	Male	32	1	1.79	0.0840	0.0818
15	Treated	Female	23	0	-1.18	0.1416	0.0420
16	Treated	Female	32	0	-1.36	0.0926	0.0381
39	Treated	Female	69	0	-2.17	0.0314	0.0690



case 1: younger male: moderate Hat, better than predicted → large Cook D

case 39: older female: small Hat, but did not improve with treatment

Looking ahead

- Logistic regression models need not always have linear effects— models **nonlinear** in Xs sometimes useful
- **Polytomous** outcomes can be handled as well
 - e.g., Improved = {"None", "Some", "Marked"}
- If ordinal,
 - the **proportional odds** model is a simple extension
 - **nested dichotomies** provides an alternative approach
- Otherwise, **multinomial logistic regression** is the way

Summary

- Model-based methods provide hypothesis tests, CIs & tests for individual terms
- Logistic regression: A `glm()` for a binary response
 - linear model for the log odds $\Pr(Y=1)$
 - All similar to classical ANOVA, regression models