






# GLMs for Count Data

Michael Friendly

Psych 6136

November 22, 2017



# Outline

- 1 Generalized linear models
- 2 GLMs for count data
  - Example: phdpubs
- 3 Model diagnostics
  - Interactions
  - Nonlinearity
  - Outliers, leverage and influence
- 4 Overdispersion
  - Quasi-poisson models
  - Negative-binomial models
- 5 Excess zeros
  - Zero-inflated models
  - Hurdle models
  - Example

# Generalized linear models

We have used generalized linear models (`glm()`) in two contexts so far:

## Loglinear models

- the outcome variable is the **vector of frequencies**  $\mathbf{y}$  in a table cross-classified by factors in a design matrix  $\mathbf{X}$
- The model is expressed as a linear model for  $\log \mathbf{y}$

$$\log(\mathbf{y}) = \mathbf{X}\beta$$

- The random (or unexplained) variation is expressed as a Poisson distribution for  $\mathcal{E}(\mathbf{y} | \mathbf{X})$

# Generalized linear models

## Logistic regression

- the outcome variable is a **categorical response**  $\mathbf{y}$ , with predictors  $\mathbf{X}$
- The model is expressed as a linear model for the log odds that  $y = 1$  vs.  $y = 0$ .

$$\text{logit}(\mathbf{y}) \equiv \log \left[ \frac{\Pr(y = 1)}{\Pr(y = 0)} \right] = \mathbf{X}\beta$$

- The random (or unexplained) variation is expressed as a Binomial distribution for  $\mathcal{E}(\mathbf{y} | \mathbf{X})$

Hey, aren't these both very like the familiar, classical linear model,

$$\mathbf{y} = \mathbf{X}\beta + \epsilon, \quad \epsilon \sim \mathcal{N}(0, \sigma^2 \mathbf{I}) \quad ?$$

Yes, for some transformation,  $g(\mathbf{y})$ , and with different distributions!

# Generalized linear models

Nelder & Wedderburn (1972) said, “Let there be light!”, a **generalized linear model**, encompassing them all, and many more. This has 3 components:

- A **random component**, specifying the conditional distribution of  $\mathbf{y}$  given the explanatory variables in  $\mathbf{X}$ , with mean  $\mathcal{E}(y_i | \mathbf{x}_i) = \mu_i$ 
  - The normal (Gaussian), binomial, and Poisson are already familiar
  - But, these are all members of an **exponential family**
  - GLMs now include an even wider family: negative-binomial and others
- The **systematic component**, a linear function of the predictors called the **linear predictor**

$$\boldsymbol{\eta} = \mathbf{X}\boldsymbol{\beta} \quad \text{or} \quad \eta_i = \beta_0 + \beta_1 X_{i1} + \cdots + \beta_p X_{ip}$$

- An invertible **link function**,  $g(\mu_i) = \eta_i = \mathbf{x}_i^T \boldsymbol{\beta}$  that transforms the expected value of the response to the linear predictor
  - The link function is invertable, so we can go back to the **mean function**  $g^{-1}(\eta_i) = \mu_i$

# Mean functions

## Standard GLM link functions and their inverses:

**Table 11.1:** Common link functions and their inverses used in generalized linear models

Link name	Function: $\eta_i = g(\mu_i)$	Inverse: $\mu_i = g^{-1}(\eta_i)$
identity	$\mu_i$	$\eta_i$
square-root	$\sqrt{\mu_i}$	$\eta_i^2$
log	$\log_e(\mu_i)$	$\exp(\eta_i)$
inverse	$\mu_i^{-1}$	$\eta_i^{-1}$
inverse-square	$\mu_i^{-2}$	$\eta_i^{-1/2}$
logit	$\log_e \frac{\mu_i}{1-\mu_i}$	$\frac{1}{1+\exp(-\eta_i)}$
probit	$\Phi^{-1}(\mu_i)$	$\Phi(\eta_i)$
log-log	$-\log_e[-\log_e(\mu_i)]$	$\exp[-\exp(-\eta_i)]$
comp. log-log	$\log_e[-\log_e(1-\mu_i)]$	$1 - \exp[-\exp(\eta_i)]$

- The top section recognizes standard transformations often used with traditional linear models
- The bottom section is for binomial data, where  $y_i$  represents an observed proportion in  $n_i$  trials

# Canonical links and variance functions

- For every distribution family, there is a default, **canonical** link function
- Each one also specifies the expected relationship between mean and variance

**Table 11.2:** Common distributions in the exponential family used with generalized linear models and their canonical link and variance functions

Family	Notation	Canonical link	Range of $y$	Variance function, $\mathcal{V}(\mu   \eta)$
Gaussian	$N(\mu, \sigma^2)$	identity: $\mu$	$(-\infty, +\infty)$	$\phi$
Poisson	$\text{Pois}(\mu)$	$\log_e(\mu)$	$0, 1, \dots, \infty$	$\mu$
Negative-Binomial	$\text{NBin}(\mu, \theta)$	$\log_e(\mu)$	$0, 1, \dots, \infty$	$\mu + \mu^2/\theta$
Binomial	$\text{Bin}(n, \mu)/n$	$\text{logit}(\mu)$	$\{0, 1, \dots, n\}/n$	$\mu(1 - \mu)/n$
Gamma	$G(\mu, \nu)$	$\mu^{-1}$	$(0, +\infty)$	$\phi\mu^2$
Inverse-Gaussian	$IG(\mu, \nu)$	$\mu^2$	$(0, +\infty)$	$\phi\mu^3$

# Variance functions and over-dispersion

- In the classical Gaussian linear model, the conditional variance is constant,  $\phi = \sigma_\epsilon^2$ .
- For binomial data, the variance function is  $\mathcal{V}(\mu_i) = \mu_i(1 - \mu_i)/n_i$ , with  $\phi$  fixed at 1
- In the Poisson family,  $\mathcal{V}(\mu_i) = \mu_i$  and the dispersion parameter is fixed at  $\phi = 1$ .
- In practice, it is common for count data to exhibit **overdispersion**, meaning that  $\mathcal{V}(\mu_i) > \mu_i$ .
- One way to correct for this is to allow the dispersion parameter to be estimated from the data, giving what is called the **quasi-Poisson** family, with  $\mathcal{V}(\mu_i) = \hat{\phi}\mu_i$ .



# Variance functions and over-dispersion

Overdispersion often results from failures of the assumptions of the model:

- supposedly independent observations may be correlated
- the probability of an event may not be constant, or
- it may vary with unmeasured or unmodeled variables

# ML Estimation

- GLMs are fit by the method of **maximum likelihood**.
- For the Poisson distribution with mean  $\mu$ , the probability that the random variable  $Y$  takes values  $y = 0, 1, 2, \dots$  is

$$\Pr(Y = y) = \frac{e^{-\mu} \mu^y}{y!}$$

- In the GLM with a log link, the mean,  $\mu_i$  depends on the predictors in  $\mathbf{x}$  through

$$\log_e(\mu_i) = \mathbf{x}_i^T \boldsymbol{\beta}$$

- The log-likelihood function (ignoring a constant) for  $n$  independent observations has the form

$$\log_e \mathcal{L}(\boldsymbol{\beta}) = \sum_{i=1}^n \{y_i \log_e(\mu_i) - \mu_i\}$$

- It can be shown that the maximum likelihood estimators are solutions to the **estimating equations**,

$$\mathbf{X}^T \mathbf{y} = \mathbf{X}^T \boldsymbol{\mu}$$

- The solutions are found by iteratively re-weighted least squares.

# Goodness of fit

- The **residual deviance** defined as twice the difference between the maximum log-likelihood for the **saturated model** that fits perfectly and maximized log-likelihood for the fitted model.

$$D(\mathbf{y}, \hat{\boldsymbol{\mu}}) \equiv 2[\log_e \mathcal{L}(\mathbf{y}; \mathbf{y}) - \log_e \mathcal{L}(\mathbf{y}; \hat{\boldsymbol{\mu}})] .$$

- For classical (Gaussian) linear models, this is just the **residual sum of squares**
- For Poisson models with a log link giving  $\boldsymbol{\mu} = \exp(\mathbf{x}^T \boldsymbol{\beta})$ , the deviance takes the form

$$D(\mathbf{y}, \hat{\boldsymbol{\mu}}) = 2 \sum_{i=1}^n \left[ y_i \log_e \left( \frac{y_i}{\hat{\mu}_i} \right) - (y_i - \hat{\mu}_i) \right] .$$

- For a GLM with  $p$  parameters, both the Pearson and residual deviance statistics follow approximate  $\chi^2_{n-p}$  distributions with  $n - p$  degrees of freedom.

# GLMs for count data

- Typically, these are fit using: `glm( y ~ x1 + x2 + x3, family=poisson, data=mydata)`
- As in other linear models, the predictors  $x_j$  can be discrete factors, quantitative variables, and so forth.
- This fixes the dispersion parameter  $\phi$  to 1, assuming that the count variable  $y$  conditional on  $x_1, x_2, \dots$  is Poisson distributed.
- It is possible to fit a **quasi Poisson** model, allowing  $\phi$  to be estimated from the data. Specify: `family=quasipoisson`. This allows the variance to be proportional to the mean,

$$\mathcal{V}(y_i | \eta_i) = \phi \mu_i$$

- Another possibility is the negative-binomial model, which has

$$\mathcal{V}(y_i | \eta_i) = \mu_i + \mu_i^2 / \theta$$

## Example: Publications of PhD Candidates

Example 3.24 in DDAR gives data on the number of publications by PhD candidates in biochemistry in the last 3 years of study

```
data("PhdPubs", package = "vcdExtra")  
table(PhdPubs$articles)
```

```
##  
##      0      1      2      3      4      5      6      7      8      9     10     11     12     16     19  
## 275 246 178  84  67  27  17  12   1   2   1   1   2   1   1
```

- Predictors are: gender, marital status, number of young children, prestige of the doctoral department, and number of publications by the student's mentor.

## Example: Publications of PhD Candidates

- Initially, ignore the predictors.
- For the Poisson, equivalent to an intercept-only model:

```
glm(articles ~ 1, data=PhdPubs, family="poisson")
```

As a quick check on the Poisson assumption:

```
with(PhdPubs, c(mean=mean(articles),  
                 var=var(articles),  
                 ratio=var(articles)/mean(articles)))
```

```
##      mean      var  ratio  
## 1.6929 3.7097 2.1914
```


The assumption that mean = variance could be met when we add predictors.

# Example: Publications of PhD Candidates


First, look at rootograms:

```
plot(goodfit(PhdPubs$articles), xlab = "Number of Articles",  
     main = "Poisson")  
plot(goodfit(PhdPubs$articles, type = "nbinomial"),  
     xlab = "Number of Articles", main = "Negative binomial")
```

Poisson



Negative binomial



One reason the Poisson doesn't fit: excess 0s (some never published?)

# Fitting the Poisson model

Fit the model with all main effects:

```
# predictors: female, married, kid5, phdprestige, mentor
phd.pois <- glm(articles ~ ., data=PhdPubs, family=poisson)
Anova(phd.pois)

## Analysis of Deviance Table (Type II tests)
##
## Response: articles
##          LR Chisq Df Pr(>Chisq)
## female      17.1  1  3.6e-05 ***
## married       6.6  1   0.01 *
## kid5         22.1  1  2.6e-06 ***
## phdprestige   1.0  1   0.32
## mentor     126.8  1 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Only `phdprestige` is NS; it does no harm to keep it, for now.



## Interpreting coefficients

$\beta_j$  is the **increment** in log (articles) for a 1 unit change in  $x_j$ ;  $\exp(\beta_j)$  is the **multiple** of articles:

```
round(cbind(beta = coef(phd.pois),
             expbeta = exp(coef(phd.pois)),
             pct = 100 * (exp(coef(phd.pois)) - 1)), 3)
```

##		beta	expbeta	pct
##	(Intercept)	0.266	1.304	30.425
##	female1	-0.224	0.799	-20.102
##	married1	0.157	1.170	17.037
##	kid5	-0.185	0.831	-16.882
##	phdprestige	0.025	1.026	2.570
##	mentor	0.025	1.026	2.555

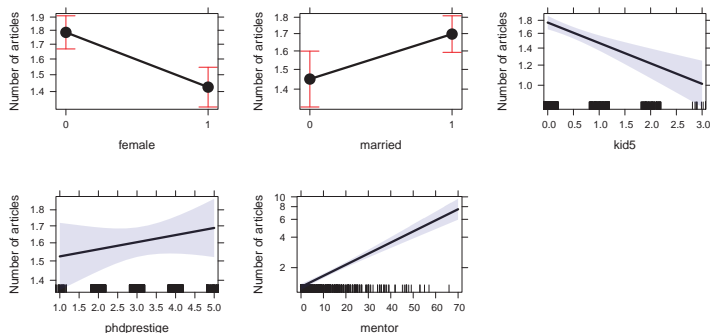
Thus:

- females publish -0.224 fewer log (articles), or  $0.8 \times$  that of males
- married publish 0.157 more log (articles); or  $1.17 \times$  unmarried (17% increase)
- each additional young child decreases this by 0.185; or  $0.831 \times$  articles (16.9% decrease)
- each mentor pub multiplies student pub by 1.026, a 2.6% increase

## Effect plots

As usual, we can understand the fitted model from predicted values for the model effects:

```
library(effects); plot(allEffects(phd.pois))
```



These are better visual summaries for a model than a table of coefficients.

# Model diagnostics

Diagnostic tests for count data GLMs are similar to those used for classical linear models

- Test for presence of interactions
  - Fit model(s) with some or all two-way interactions
- Non-linear effects of quantitative predictors?
  - Component-plus-residual plots— `car::crPlot()` are useful here
- Outliers? Influential observations?
  - `car::influencePlot()` is your friend

For count data models, we should also check for **over-dispersion**. This is similar to **homogeneity of variance** checks in `lm()`

# Testing for interactions

As a quick check for interactions, fit the model with all two-way terms

```
phd.pois1 <- update(phd.pois, . ~ .^2)
Anova(phd.pois1)
```

```
## Analysis of Deviance Table (Type II tests)
##
## Response: articles
##
##          LR Chisq Df Pr(>Chisq)
## female          14.5  1    0.00014 ***
## married           6.2  1    0.01277 *
## kid5             19.5  1    9.8e-06 ***
## phdprestige       1.0  1    0.32655
## mentor          128.1  1    < 2e-16 ***
## female:married     0.3  1    0.60995
## female:kid5        0.1  1    0.72929
## female:phdprestige  0.2  1    0.63574
## female:mentor      0.0  1    0.91260
## married:kid5              0
## married:phdprestige  1.7  1    0.19153
## married:mentor      1.2  1    0.28203
## kid5:phdprestige     0.2  1    0.68523
## kid5:mentor         2.8  1    0.09290 .
## phdprestige:mentor   3.8  1    0.05094 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

# Compare models I

Compare models: LR tests for nested models (**anova()**), and AIC/BIC (**LRstats()**)

```
anova(phd.pois, phd.pois1, test="Chisq")

## Analysis of Deviance Table
##
## Model 1: articles ~ female + married + kid5 + phdprestige + mentor
## Model 2: articles ~ female + married + kid5 + phdprestige + mentor + fe
##      female:kid5 + female:phdprestige + female:mentor + married:kid5 +
##      married:phdprestige + married:mentor + kid5:phdprestige +
##      kid5:mentor + phdprestige:mentor
##      Resid. Df Resid. Dev Df Deviance Pr(>Chi)
## 1           909         1634
## 2           900         1618   9      15.2    0.086 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

## Compare models II

```
LRstats(phd.pois, phd.pois1)



## Likelihood summary table:
##           AIC   BIC LR Chisq  Df Pr(>Chisq)
## phd.pois  3313 3342    1634 909    <2e-16 ***
## phd.pois1 3316 3388    1618 900    <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

- There seems to be no reason to include interactions in the model
- We might want to re-visit this, after examining other models for the basic count distribution (quasi-poisson, negative-binomial)

# Basic model plots

Only two of the standard model plots are informative for count data models

```
plot(phd.pois, which=c(1,5))
```



# Nonlinearity diagnostics


- Non-linear relations are difficult to assess in **marginal plots**, because they don't control (or adjust) for other predictors
- **Component-plus-residual** plots (also called **partial residual plots**) can show non-linear relations for numeric predictors
  - These graph the value of  $\hat{\beta}_i x_i + \text{residual}_i$  vs. the predictor,  $x_i$ .
  - In this plot, the slope of the points is the coefficient,  $\hat{\beta}_i$  in the full model
  - The residual is  $y_i - \hat{y}_i$  in the full model
- A non-parametric (e.g., **loess()**) smooth makes it easy to detect non-linearity



# Nonlinearity diagnostics: `car::crPlot()`

Is the relationship between articles published by the student and the mentor adequately represented as linear?

```
crPlot(phd.pois, "mentor", pch=16, lwd=4, id.n=2)
```



# Residuals I

Several types of residuals can be defined based on the Pearson and deviance goodness-of-fit measures

- the **Pearson residual** is the case-wise contribution to Pearson  $\chi^2$

$$r_i^P = \frac{y_i - \hat{\mu}_i}{\sqrt{\hat{v}(y_i)}}$$

- the **deviance residual** is the signed square root of the contribution to the deviance  $G^2$

$$r_i^D = \text{sign}(y_i - \hat{\mu}_i) \sqrt{d_i}$$

- Both of these have **standardized** forms that correct for conditional variance and leverage, and have approx.  $\mathcal{N}(0, 1)$  distributions.

$$\tilde{r}_i^P = \frac{r_i^P}{\sqrt{\hat{\phi}(1 - h_i)}}$$


$$\tilde{r}_i^D = \frac{r_i^D}{\sqrt{\hat{\phi}(1 - h_i)}}$$

## Residuals II

- The most useful is the ***studentized residual*** (or deletion residual), `rstudent()` in R. This estimates the standardized residual resulting from omitting each observation in turn. An approximation is:

$$\tilde{r}_i^S = \text{sign}(y_i - \hat{\mu}_i) \sqrt{(1 - h_i)(\tilde{r}_i^D)^2 + h_i(\tilde{r}_i^P)^2} .$$

# Outliers, leverage and influence



`influencePlot(phd.pois)`

- Several observations (913–915) stand out with large + residuals
- One observation (328) has a large leverage
- Why are they unusual? Do they affect our conclusions?
- Look back at data & decide what to do!

# Outliers, leverage and influence

At the very least, we should look at these observations in the data:

```
PhdPubs[c(328, 913:915), ]
```

##	articles	female	married	kid5	phdprestige	mentor
## 328	1	0	1	1	2	77
## 913	12	0	1	1	2	5
## 914	16	0	1	0	2	21
## 915	19	0	1	0	2	42

- case 328: Mentor published 77 papers! Student, only 1
- 913–915: all published >> predicted

# Outlier test

- A formal test for outliers can be based on the **studentized residuals**, `rstudent(model)`, using the standard normal distribution for  $p$ -values
- A Bonferroni correction should be applied, because interest focuses on the **largest**  $n$  absolute residuals.

For the Poisson model, 4 observations are nominated as large + outliers:

```
outlierTest(phd.pois, cutoff=0.001)
```

##		rstudent	unadjusted p-value	Bonferonni p
##	914	5.5423	2.9852e-08	2.7315e-05
##	913	5.3821	7.3617e-08	6.7360e-05
##	911	5.2074	1.9153e-07	1.7525e-04
##	915	5.1504	2.5988e-07	2.3779e-04

# Overdispersion

- The Poisson model for counts assumes  $\mathcal{V}(\mu_i) = \mu_i$ , i.e., the dispersion parameter  $\phi = 1$
- But often, the counts exhibit greater variance than the Poisson distribution allows,  $\mathcal{V}(\mu_i) > \mu_i$  or  $\phi > 1$ 
  - The observations (counts) may not be independent (clustering)
  - The probability of an “event” may not be constant
  - There may be unmeasured influences, not accounted for in the model
  - These effects are sometimes called “unmodeled heterogeneity”
- The consequences are:
  - Standard errors of the coefficients,  $\text{se}(\hat{\beta}_j)$  are optimistically small
  - Wald tests,  $z_j = \hat{\beta}_j / \text{se}(\hat{\beta}_j)$ , are too large, and thus overly liberal.

# Testing overdispersion

- Statistical tests for overdispersion are described in DDAR §11.3.4.
- They test  $H_0 : \mathcal{V}(y) = \mu$ , vs.  $H_1$  that variance depends on the mean according to some function  $f(\mu)$

$$\mathcal{V}(y) = \mu + \alpha \times f(\mu)$$

- This is implemented in `dispersiontest()` in the **AER** package.
  - If significant, overdispersion should not be ignored
  - Alternatively, you can try fitting a more general model to see what difference it makes.



# Overdispersion: Quasi-poisson models

- Instead, we can fit another version of the model in which the dispersion  $\phi$  is a free parameter, estimated along with the other coefficients. That is, the conditional variance is allowed to be

$$\mathcal{V}(y_i | \eta_i) = \phi \mu_i$$

- This model is fit with `glm()` using `family=quasipoisson`
  - the estimated coefficients  $\hat{\beta}$  are **unchanged**
  - the standard errors are multiplied by  $\hat{\phi}^{1/2}$
  - peace, order, and good governance is restored!

# Overdispersion: Quasi-poisson models

- One estimate of the dispersion parameter is the residual deviance divided by degrees of freedom  $\hat{\phi} = D(\mathbf{y}, \hat{\mu}) / df$
- The Pearson  $\chi^2$  statistic has better statistical properties and is more commonly used

$$\hat{\phi} = \frac{X_P^2}{n - p} = \sum_{i=1}^n \frac{(y_i - \hat{\mu}_i)^2}{\hat{\mu}_i} / (n - p) \quad .$$

For the PhdPubs data, these estimates are quite similar: about 80% overdispersion

```
with(phd.pois, deviance / df.residual)

## [1] 1.7971

sum(residuals(phd.pois, type = "pearson")^2) / phd.pois$df.residual

## [1] 1.8304
```

# Fitting the quasi-poisson model

The quasi-Poisson model is can be fit using `glm()` as:

```
phd.qpois <- glm(articles ~ ., data=PhdPubs, family=quasipoisson)
```

The dispersion parameter estimate  $\hat{\phi}$  can be obtained as follows:

```
(phi <- summary(phd.qpois)$dispersion)
```

```
## [1] 1.8304
```

This is much better than variance/mean ratio of 2.91 calculated for the marginal distribution ignoring the predictors.

Coefficients unchanged; std. errors multiplied by  $\hat{\phi}^{1/2} = \sqrt{1.83} = 1.35$ .

```
summary(phd.gpois)
```

```
##
## Call:
## glm(formula = articles ~ ., family = quasipoisson, data = PhdPubs)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -3.488  -1.538  -0.365   0.577   5.483
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  0.26562    0.13478   1.97  0.04906 *
## female1     -0.22442    0.07384  -3.04  0.00244 **
## married1     0.15732    0.08287   1.90  0.05795 .
## kid5        -0.18491    0.05427  -3.41  0.00069 ***
## phdprestige  0.02538    0.03419   0.74  0.45815
## mentor      0.02523    0.00275   9.19 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for quasipoisson family taken to be 1.8304)
##
##      Null deviance: 1817.4  on 914  degrees of freedom
## Residual deviance: 1633.6  on 909  degrees of freedom
## AIC: NA
##
## Number of Fisher Scoring iterations: 5
```

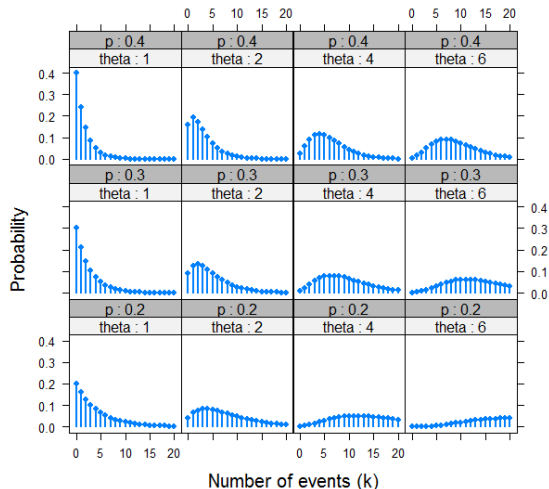
# The negative-binomial model

- The negative-binomial model is a different generalization of the Poisson that allows for over-dispersion
- Mathematically, it allows the mean  $\mu | \mathbf{x}_i$  to vary across observations as a gamma distribution with a shape parameter  $\theta$ .
- The variance function,  $\mathcal{V}(y_i) = \mu_i + \mu_i^2/\theta$ , allows the variance of  $y$  to increase more rapidly than the mean.
- Another parameterization uses  $\alpha = 1/\theta$

$$\mathcal{V}(y_i) = \mu_i + \mu_i^2/\theta = \mu_i + \alpha\mu_i^2 ,$$

- As  $\alpha \rightarrow 0$ ,  $\mathcal{V}(y_i) \rightarrow \mu_i$  and the negative-binomial converges to the Poisson.

# The negative-binomial model



Overdispersion decreases as  $\theta$  increases


# The negative-binomial model: Fitting

- For fixed  $\theta$ , the negative-binomial is another special case of the GLM
- This is handled in the **MASS** package, with `family=negative.binomial(theta)`
- But most often,  $\theta$  is unknown, and must be estimated from the data
- This is implemented in `glm.nb()` in the **MASS** package.

```
library(MASS)
phd.nbin <- glm.nb(articles ~ ., data=PhdPubs)
```

# Visualizing the mean variance relation

One way to see the difference among models is to plot the variance vs. mean for **grouped** values of the fitted linear predictor.





- The smoothed (loess) curve gives the **empirical mean-variance** relationship
- Also plot the theoretical mean-variance from different models
- For PhdPubs, the data is most similar to the negative-binomial
- The models differ most for those with  $> 3$  articles



# Visualizing goodness-of-fit

The `countreg` package extends the `rootogram()` function to work with fitted models:

```
countreg::rootogram(phd.pois, main="PhDPubs: Poisson")  
countreg::rootogram(phd.nbin, main="PhDPubs: Negative-Binomial")
```



The Poisson model shows a systematic, wave-like pattern with excess zeros, too few observed frequencies for counts of 1–3.

# What difference does it make?

The NB is certainly a better fit than the Poisson; the QP cannot be distinguished by standard tests

```
LRstats(phd.pois, phd.qpois, phd.nbin)

## Likelihood summary table:
##           AIC   BIC LR Chisq  Df Pr(>Chisq)
## phd.pois  3313 3342    1634  909    <2e-16 ***
## phd.qpois                909
## phd.nbin   3135 3169    1004  909    0.015 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Can also compare standard errors of the coefficients:

```
##           pois qpois  nbin
## (Intercept) 0.100 0.135 0.133
## female1     0.055 0.074 0.073
## married1    0.061 0.083 0.082
## kid5         0.040 0.054 0.053
## phdprestige  0.025 0.034 0.034
## mentor      0.002 0.003 0.003
```

# What have we learned?

A summary for an article to this point would use the result of negative-binomial model, from `summary(phd.nbin)`

- The number of articles published by these PhD candidates is most strongly influenced by publications of their mentor
- Increasing young children (`kids5`) results in fewer publications.
- Being married is marginally non-significant— don't interpret
- The prestige of the university doesn't make a difference
- There are still some remaining doubts:
  - Several cases (328, 913–915) appeared unusual in earlier diagnostic plots. Refit without them to see if any conclusions change.
  - The NB model seems to account for the zero counts— students who never published.
  - Is there a better way?

# Excess zero counts


- A common problem in count data models is that many sets of data have more observed zero counts than the (quasi) Poisson or NB models can handle.
  - In the PhdPubs data, 275 of 915 (30%) candidates published zilch, bupkis
  - The expected count of 0 articles in the Poisson model is only 191 (21%)
- Maybe there are two types of students giving zero counts:
  - Those who never intend to publish (non-academic career path?)
  - The rest, who do intend to publish, but have not yet done so
  - This suggests the idea of **zero inflation**
- An alternative idea is that there is some **hurdle** to overcome before attaining a positive count, e.g., external pressure from the mentor.

Beyond simply identifying this as a problem of lack-of-fit, understanding the **reasons** for excess zero counts can contribute to a more complete explanation of the phenomenon of interest.


## Two model types for excess zeros

- **zero-inflated models:** The responses with  $y_i = 0$  arise from a mixture of **structural**, always 0 values, with  $\Pr(y_i = 0) = \pi_i$  and the rest, which are **random** 0s, with  $\Pr(y_i = 0) = 1 - \pi_i$
- **hurdle models:** One process determines whether  $y_i = 0$  with  $\Pr(y_i = 0) = \pi_i$ . A second process determines the distribution of values of positive counts,  $\Pr(y_i | y_i > 0)$

Zero-inflated



Hurdle



# Zero-inflated models

The **zero-inflated Poisson** (ZIP) model has two components:

- A logistic regression model for membership in the unobserved (latent) class of those for whom  $y_i$  is necessarily zero

$$\text{logit}(\pi_i) = \mathbf{z}_i^T \boldsymbol{\gamma} = \gamma_0 + \gamma_1 z_{i1} + \gamma_2 z_{i2} + \cdots + \gamma_q z_{iq} .$$

- A Poisson model for the other class (e.g., “publishers”), for whom  $y_i$  may be 0 or positive.

$$\log_e \mu(y_i | \mathbf{x}_i) = \mathbf{x}_i^T \boldsymbol{\beta} = \beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \cdots + \beta_q x_{ip} .$$

In applications, the same predictors can be (and often are) used in both models ( $\mathbf{x} = \mathbf{z}$ ).

## Zero-inflated models

In the ZIP model, the probabilities of observing counts of  $y_i = 0$  and  $y_i > 0$  are:

$$\begin{aligned}\Pr(y_i = 0 \mid \mathbf{x}, \mathbf{z}) &= \pi_i + (1 - \pi_i)e^{-\mu_i} \\ \Pr(y_i \mid \mathbf{x}, \mathbf{z}) &= (1 - \pi_i) \times \left[ \frac{\mu_i^{y_i} e^{-\mu_i}}{y_i!} \right], \quad y_i \geq 0 .\end{aligned}$$

The conditional expectation and variance of  $y_i$  then are:

$$\begin{aligned}\mathcal{E}(y_i) &= (1 - \pi_i) \mu_i \\ \mathcal{V}(y_i) &= (1 - \pi_i) \mu_i (1 + \mu_i \pi_i) .\end{aligned}$$

When  $\pi_i > 0$ , the mean of  $y$  is always less than  $\mu_i$ ; the variance of  $y$  is greater than its mean by a dispersion factor of  $(1 + \mu_i \pi_i)$ .

The model for the count variable could also be negative-binomial, giving a *zero-inflated negative-binomial* (ZINB) model using  $\text{NBin}(\mu, \theta)$

## Zero-inflated data

Generate some random data from  $\text{Pois}(3) = \text{ZIP}(3, \pi = 0)$  and  $\text{ZIP}(3, \pi = 0.3)$ . This uses `rzipois()` in the `VGAM`.

```
library(VGAM)
set.seed(1234)
data1 <- rzipois(200, 3, 0)
data2 <- rzipois(200, 3, .3)
```

Tables of the counts:

```
table(data1)
```

```
## data1
##  0  1  2  3  4  5  6  7  8  9
## 10 31 46 54 24 20  7  3  4  1
```



```
table(data2)
```

```
## data2
##  0  1  2  3  4  5  6  7  9
## 62 26 33 31 22  9  8  8  1
```



# Zero-inflated data

Bar plots of the counts:



The 30% extra zeros decrease the mean and inflate the variance

# Hurdle models

The **Hurdle** model also has two components:

- A logistic regression model, for the probability that  $y_i = 0$  vs.  $y_i > 0$

$$\text{logit} \left[ \frac{\Pr(y_i = 0)}{\Pr(y_i > 0)} \right] = \mathbf{z}_i^T \boldsymbol{\gamma} = \gamma_0 + \gamma_1 z_{i1} + \gamma_2 z_{i2} + \cdots + \gamma_q z_{iq} .$$

- A model for the **positive** counts, taken as a **left-truncated** Poisson or negative-binomial, excluding the zero counts
- Comparing the ZIP and Hurdle models:
  - In ZIP models, the first (latent) process generates **extra** zeros (with probability  $\pi_i$ ).
  - In Hurdle models,  $y_i = 0$  and  $y_i > 0$  are fully observed. The first process generates all the zeros.

# Fitting ZIP and Hurdle models

In R, these models can be fit using the `pscl` and `countreg` packages.

`countreg` is more mature, but is only available on R-Forge, not on CRAN. Use:

```
install.packages("countreg", repos="http://R-Forge.R-project.org")
```

The functions have the following arguments:

```
zeroinfl(formula, data, subset, na.action, weights, offset,  
          dist = c("poisson", "negbin", "geometric", "binomial"),  
          ...)  
  
hurdle(formula, data, subset, na.action, weights, offset,  
        dist = c("poisson", "negbin", "geometric", "binomial"),  
        ...)
```


The formula,  $y \sim x_1 + x_2 + \dots$  uses the same predictors for both models.

Using  $y \sim x_1 + x_2 + \dots | z_1 + z_2 + \dots$  allows separate predictors for the 0 submodel.

## Visualizing zero counts

It is often useful to plot the data for the binary distinction between  $y_i = 0$  vs.  $y_i > 0$  as in logistic regression models.

```
plot(factor(articles==0) ~ mentor, data=PhdPubs,  
     ylevels=2:1, ylab="Zero articles",  
     breaks=quantile(mentor, probs=seq(0,1,.2)), cex.lab=1.25)
```



## Example: Phd Publications

Just to illustrate, we fit all four models, the combinations of (ZI, Hurdle)  $\times$  (Poisson, NBin) to the PhdPubs data.

For simplicity, we use all predictors for both the zero model and the non-zero model.

```
library(countreg)
phd.zip <- zeroinfl(articles ~ ., data=PhdPubs, dist="poisson")
phd.znb <- zeroinfl(articles ~ ., data=PhdPubs, dist="negbin")

phd.hp <- hurdle(articles ~ ., data=PhdPubs, dist="poisson")
phd.hnb <- hurdle(articles ~ ., data=PhdPubs, dist="negbin")
```

# Example: Phd Publications

Compare models, sorting by BIC:

```
LRstats(phd.pois, phd.nbin, phd.zip, phd.znb, phd.hp, phd.hnb,
        sortby="BIC")

## Likelihood summary table:
##           AIC   BIC LR Chisq  Df Pr(>Chisq)
## phd.pois  3313  3342      3301  909    <2e-16 ***
## phd.hp     3235  3292      3211  903    <2e-16 ***
## phd.zip    3234  3291      3210  903    <2e-16 ***
## phd.hnb    3131  3194      3105  902    <2e-16 ***
## phd.znb    3126  3188      3100  902    <2e-16 ***
## phd.nbin   3135  3169      3121  909    <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

The standard negative binomial looks best by BIC. Why do you think this is?

## Test the coefficients in the ZIP model using `lmtest::coeftest()`

```
library(lmtest)
coeftest(phd.zip)

##
## t test of coefficients:
##
##               Estimate Std. Error t value Pr(>|t|)
## count_(Intercept)  0.59918    0.11861    5.05 5.3e-07 ***
## count_female1     -0.20879    0.06353   -3.29 0.0011 **
## count_married1     0.10623    0.07097    1.50 0.1348
## count_kid5        -0.14271    0.04744   -3.01 0.0027 **
## count_phdprestige  0.00700    0.02981    0.23 0.8145
## count_mentor       0.01785    0.00233    7.65 5.3e-14 ***
## zero_(Intercept)  -0.56332    0.49405   -1.14 0.2545
## zero_female1       0.10816    0.28173    0.38 0.7011
## zero_married1     -0.35558    0.31796   -1.12 0.2637
## zero_kid5          0.21974    0.19658    1.12 0.2639
## zero_phdprestige  -0.00537    0.14118   -0.04 0.9697
## zero_mentor       -0.13313    0.04643   -2.87 0.0042 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Only `mentor` is significant for the zero model!

Re-fit the ZIP and ZNB models using only `mentor` for the zero models:

```
phd.zip1 <- zeroinfl(articles ~ . | mentor, data=PhdPubs, dist="poisson")
phd.znb1 <- zeroinfl(articles ~ . | mentor, data=PhdPubs, dist="negbin")
```

Compare again:

```
LRstats(phd.pois, phd.nbin, phd.zip, phd.znb, phd.hp, phd.hnb,
        phd.zip1, phd.znb1, sortBy="BIC")
```

```
## Likelihood summary table:
##           AIC   BIC LR Chisq  Df Pr(>Chisq)
## phd.pois  3313  3342      3301  909    <2e-16 ***
## phd.hp     3235  3292      3211  903    <2e-16 ***
## phd.zip    3234  3291      3210  903    <2e-16 ***
## phd.zip1   3227  3266      3211  907    <2e-16 ***
## phd.hnb    3131  3194      3105  902    <2e-16 ***
## phd.znb    3126  3188      3100  902    <2e-16 ***
## phd.nbin   3135  3169      3121  909    <2e-16 ***
## phd.znb1   3124  3168      3106  906    <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Now, the `phd.znb1` model is best by BIC. Why?



# Model interpretation: Coefficients

Ignoring NS coefficients in the revised ZNB model ([phd.znb1](#))

```
coef(phd.znb1) [c(1, 2, 4, 6, 7, 8)]
```

```
## count_(Intercept)      count_female1      count_kid5      count_mentor
##           0.357194      -0.211573      -0.167527      0.024057
## zero_(Intercept)      zero_mentor
##           -0.816912      -0.608024
```

- Count model:

$$\log(\text{articles}) = 0.357 - 0.21 \text{ female} - 0.17 \text{ kids5} + 0.024 \text{ mentor}$$

- Zero model:

$$\text{logit}(\text{articles} = 0) = -0.817 - 0.608 \text{ mentor}$$


Can you describe these in words?

# Model interpretation: Effect plots


- The `effects` package cannot yet handle zero-inflated or hurdle models.
- But the fitted values don't differ very much among these models
- Here, I use the `phd.nbin` model, and just show the effects for the important terms

```
plot(allEffects(phd.nbin)[c(1,3,5)], rows=1, cols=3)
```


female effect plot



kid5 effect plot




mentor effect plot



The ZIP sub-model for the zero counts (“did not publish”) can also be interpreted visually

- As an approximation, fit a separate logistic model for `articles==0`
- The effect plot for that gives an interpretation of the zero model.

```
phd.zero <- glm(articles==0 ~ mentor, data=PhdPubs, family=binomial)
plot(allEffects(phd.zero), main="Mentor effect on not publishing")
```



# What have we learned?

- The simple Poisson regression model fits very badly
  - Standard errors do not reflect overdispersion
  - Inference about model effects is compromised by **overly liberal** tests
- The quasi-poisson model corrects for overdispersion.
  - But doesn't account for excess 0s
- The negative-binomial model provides valid tests and fits the 0 counts well.
  - But it doesn't provide any insight into why there are so many 0s
- The ZIP and ZNB models fit well, and account for the 0s.
  - But they lose here on BIC (and AIC) measures, because they have  $2\times$  the number of parameters.
  - For simplicity, I have slighted the analogous hurdle models

# What have we learned?

- The revised ZNB model ([phd.znb1](#)), with only [mentor](#) predicting 0s, wins on parsimony, and has a simple interpretation.
  - The log odds that a student [does not publish](#) decrease by 0.61 for every article published by the mentor
  - Each mentor pub increases student publications by about 2.5%
  - $\Rightarrow$  Encourage or help your supervisor to publish!
  - (Or, choose a high publishing one.)
- For this data set, the main substantive interpretation and predicted effects are similar across models. But [details matter](#)!
- In data sets where there are substantive reasons for excess 0s, the ZI and hurdle models provide [different](#) explanations.
  - It is not always just a matter of model fit!
  - Hurdle models make the distinction between 0 and  $> 0$  more explicit
  - In ZI models, the interpretation of the mean count is clearer.

# What else is there?

The `PhdPubs` example was rather simple, in that:

- There were only a few predictors
  - Model selection methods could be based on simple `Anova()`s or `coeftest()`s
  - No need for more complex model selection methods, or cross-validation
- Of the quantitative predictors, only `mentor` and `kids5` had important effects
  - The effects of `mentor` and `kids5` were sufficiently linear.
  - No need to try polynomial (`poly(mentor, 2)`) or other non-linear effects
- There were no important interactions.
  - In Psychology, these are often called “moderator” effects
  - They complicate the interpretation if you rely tests of coefficients or “simple effects”
  - Interpretation is usually simplified in effect plots.

# What else is there?

- The response variable, `articles` was measured only once, i.e., there is no `longitudinal` aspect of the analysis.
  - One extension might track the number of articles published by these students over stages in their career.
  - Longitudinal models are examples of multilevel or hierarchical linear models
  - Well-developed for classical, Gaussian models (`lm()` → `lme4::lmer()` )
  - These models are being extended to GLMs for count data (e.g., `lme4::glmer()` )
- There was only one response variable: `articles`.
  - Another extension might analyse articles published and the number of job interviews upon graduation as a `multivariate GLM`
  - Yet another, could try to develop a `structural equation model` (SEM) or `path analysis model`, with a variable like “hired within one year?” as the ultimate binary outcome.